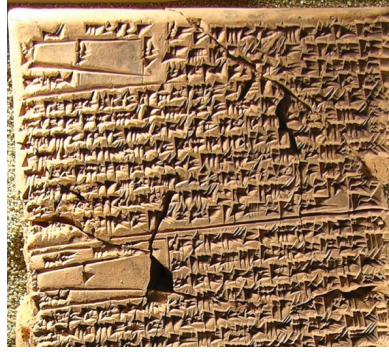


MathVista: Why Math Reasoning in Visual Contexts?



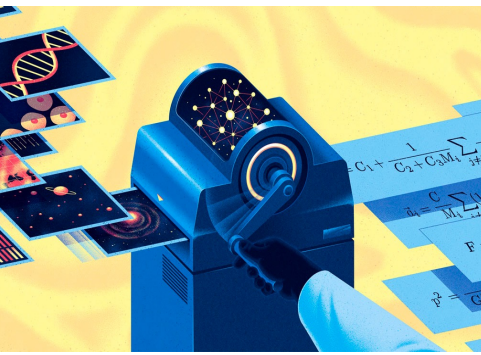
Mesopotamia, 2000BC



K-12 Math Education



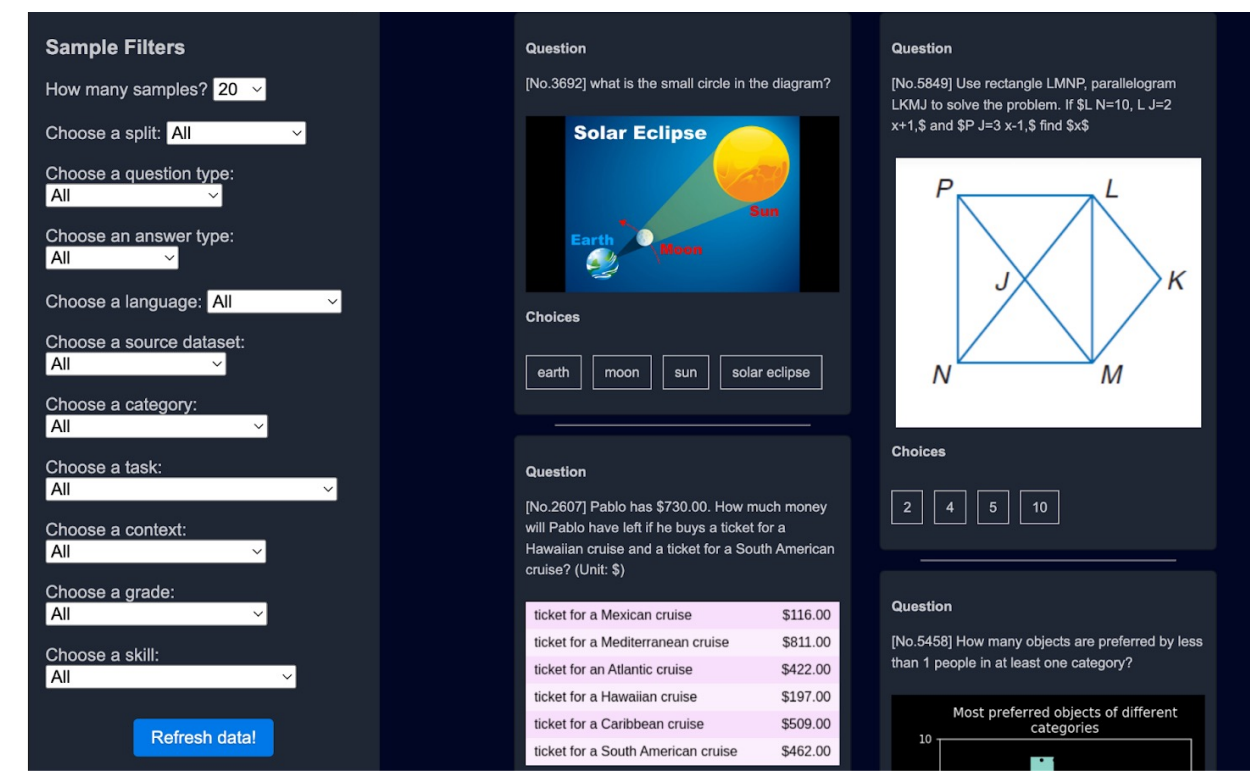
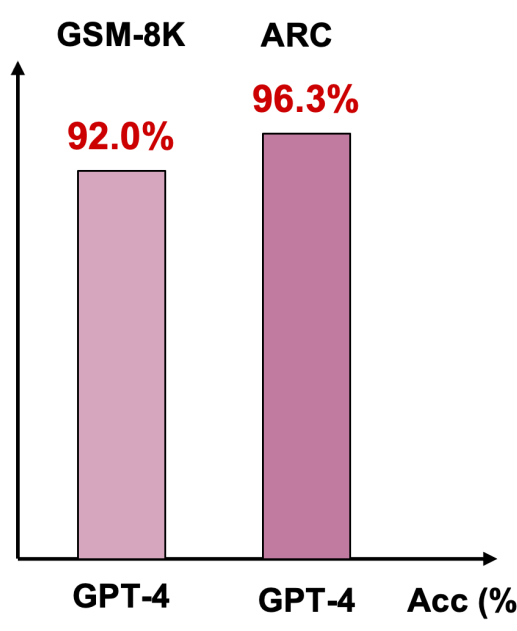
Data Analysis



Scientific Discovery

Existing Challenges

- Text-only Benchmarks
- Narrow-scope Benchmarks
- Benchmark Saturation
- Missing Systematic Evaluations

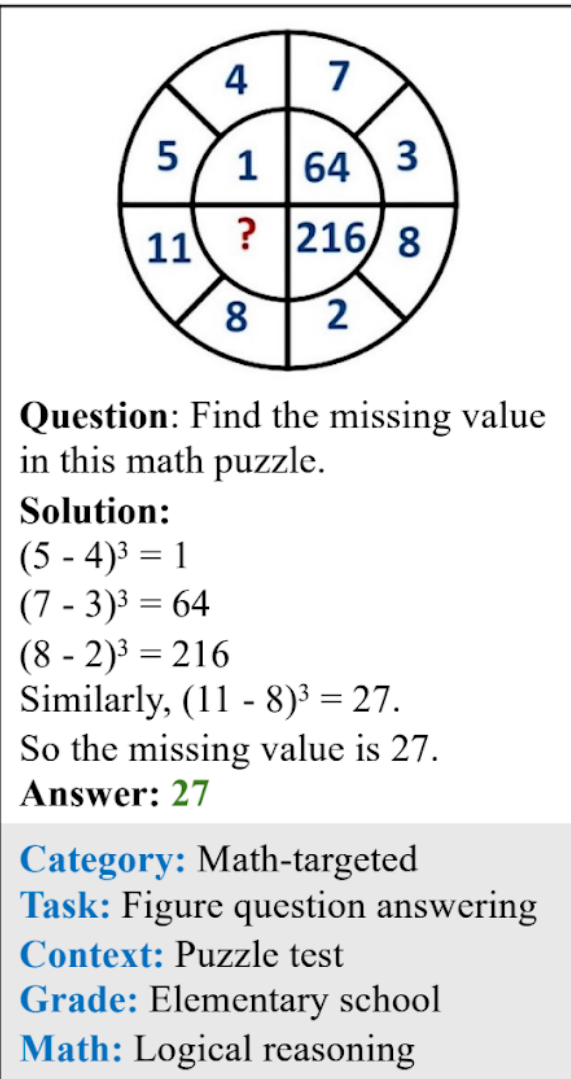


The creation of a novel testbed is **essential** to

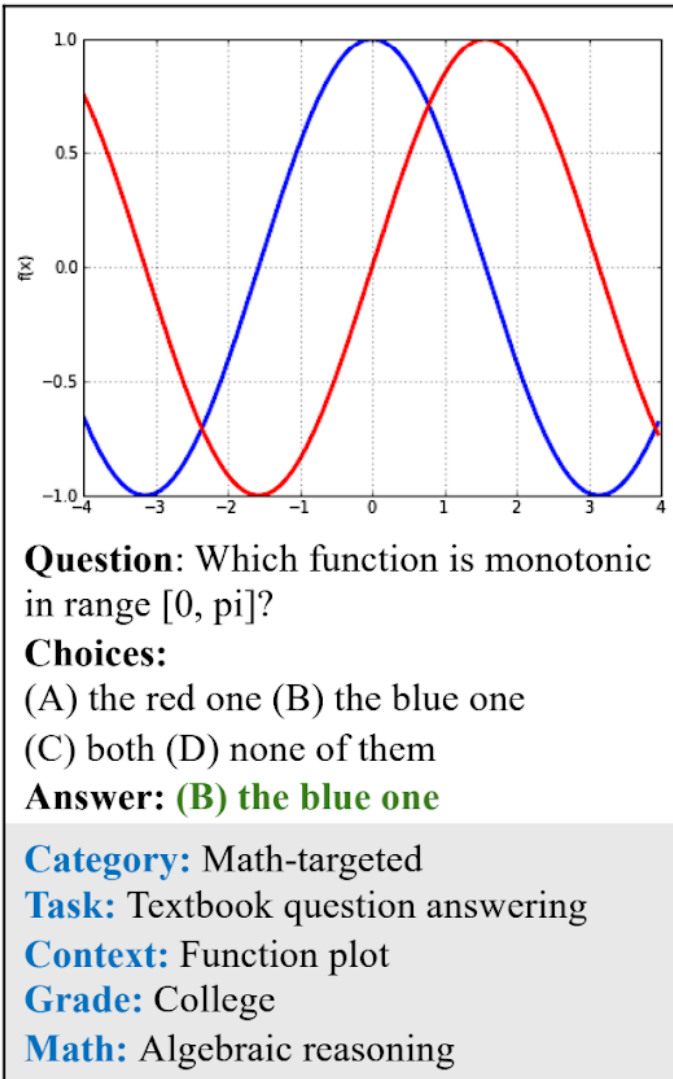
- **Facilitate** the development of **mathematical reasoning systems** in visually intensive scenarios
- **Evaluate** the research progress of **foundation models**, especially regarding their capabilities in rigorous reasoning

The MathVista Benchmark

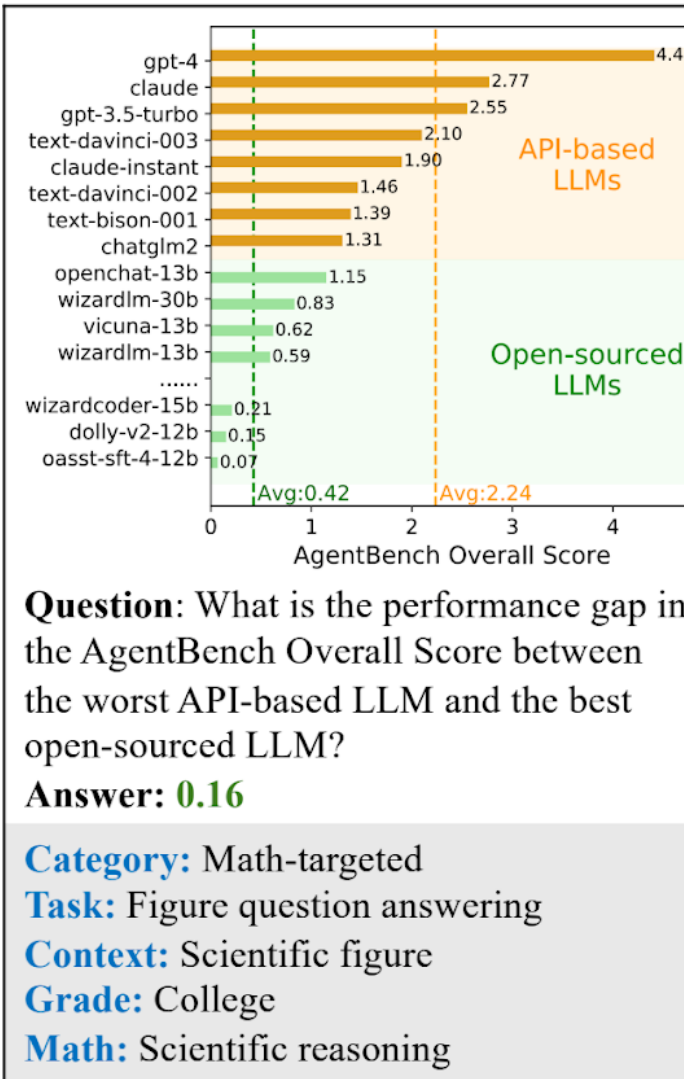
Statistic	Number
Total questions	6,141
- multiple-choice questions	3,392 (55.2%)
- Free-form questions	2,749 (44.8%)
- Questions with annotations	5,261 (85.6%)
- Questions newly annotated	736 (12.0%)
Unique number of images	5,487
Unique number of questions	4,746
Unique number of answers	1,464
Source datasets	31
- Existing VQA datasets	19
- Existing MathQA datasets	9
- Our newly annotated datasets	3
Visual context (image) classes	19
Maximum question length	213
Maximum answer length	27
Maximum choice number	8
Average question length	15.6
Average answer length	1.2
Average choice number	3.4



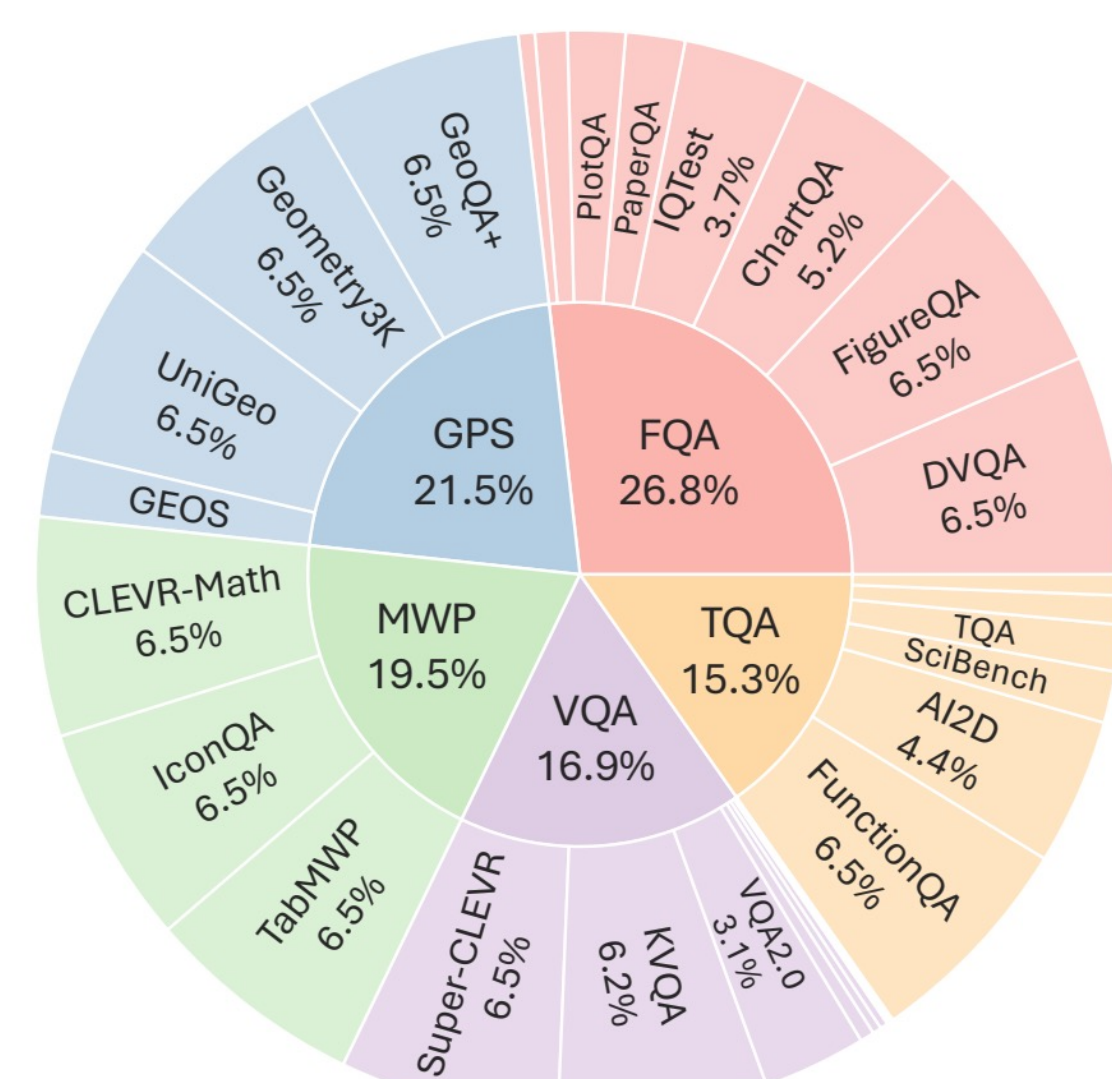
(a) IQTest



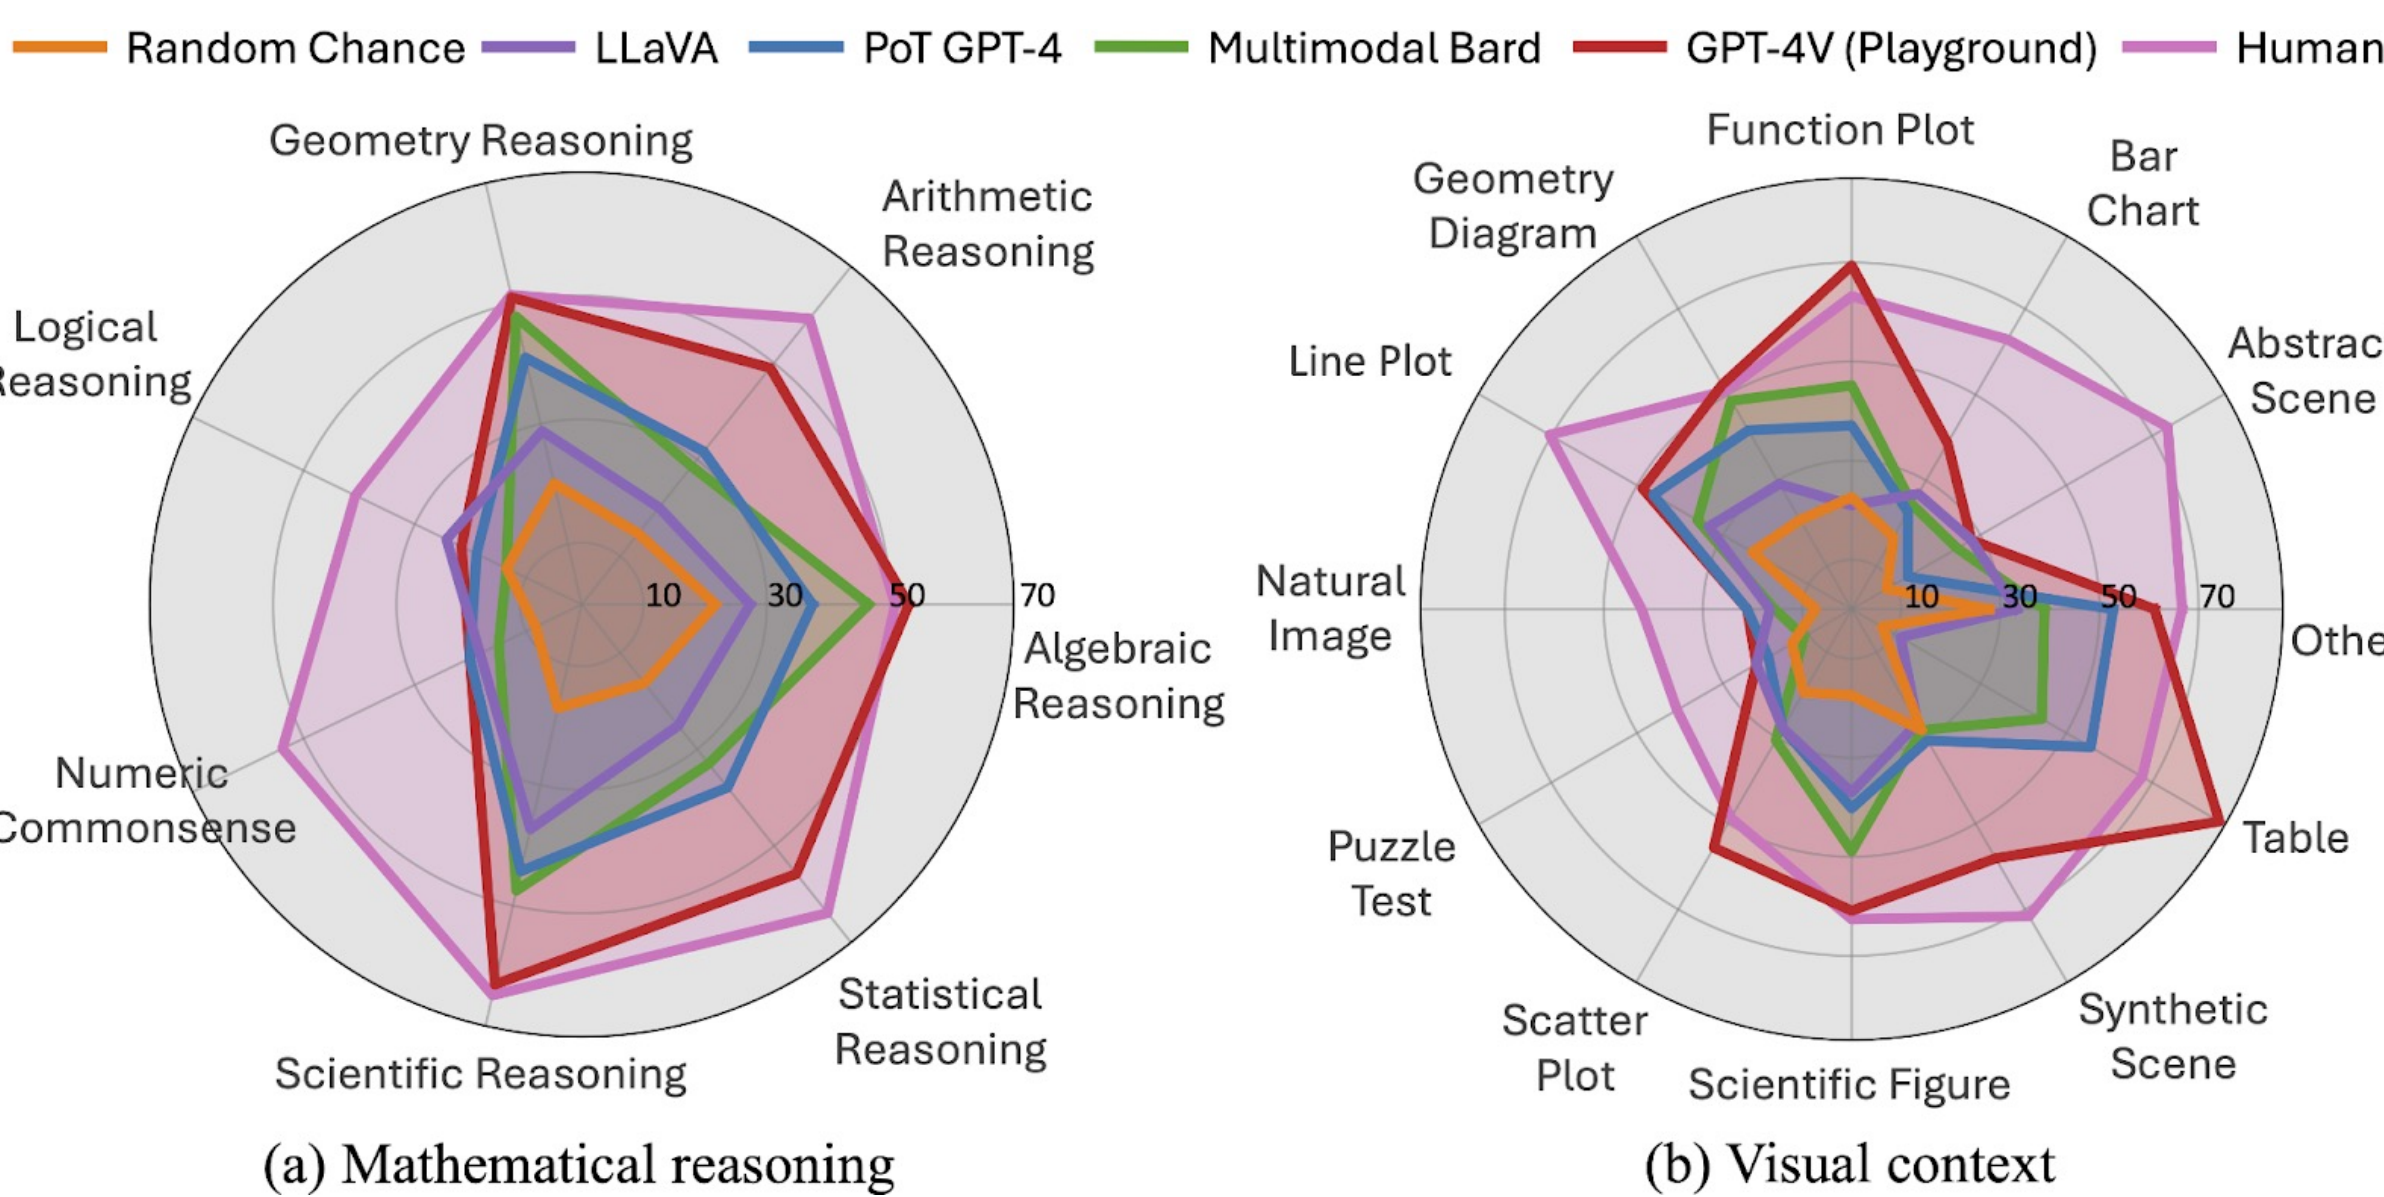
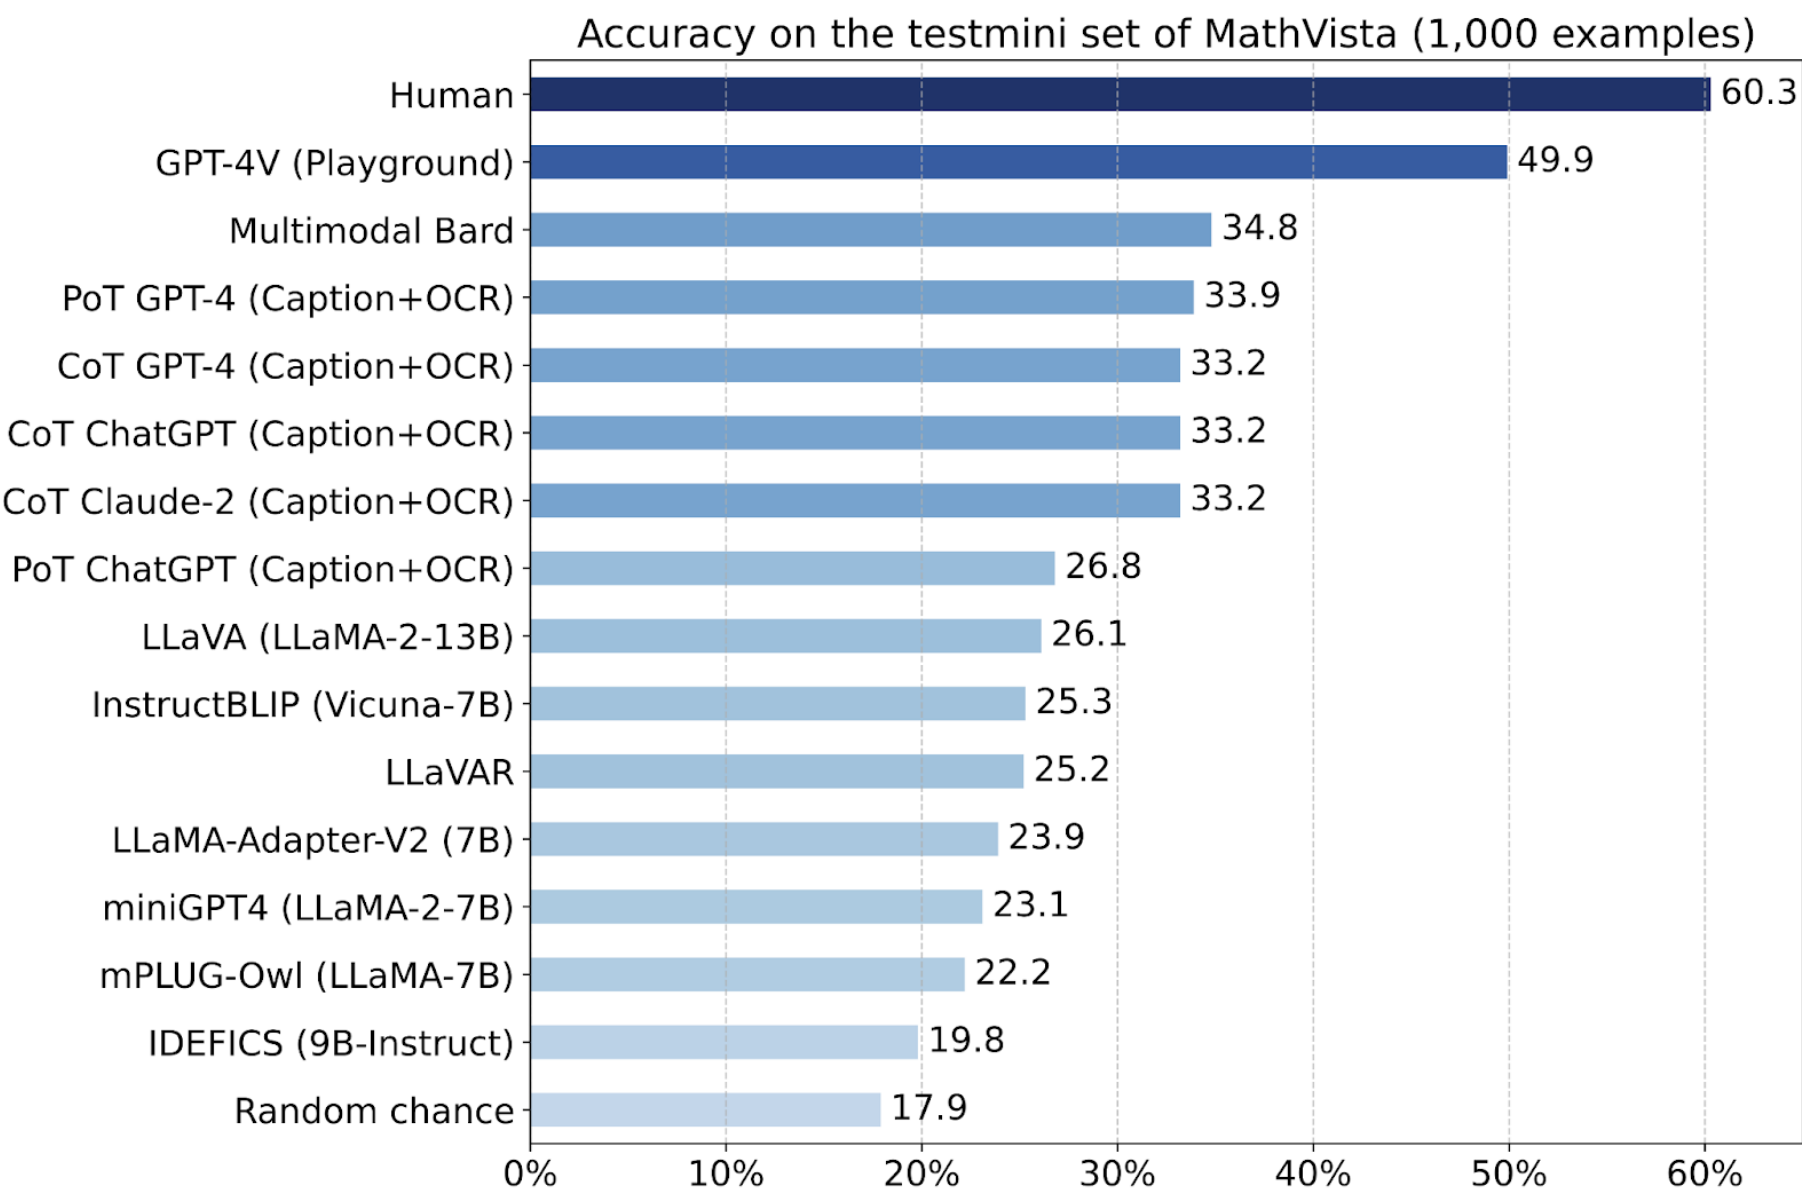
(b) FunctionQA



(c) PaperQA



Quantitative Evaluations with 12 Foundation Models (GPT-4V, Bard, and More)



Augmented LLM (GPT-4)

Laps driven	
Day	Number of laps
Thursday	53
Friday	53
Saturday	53
Sunday	55
Monday	55

Question: A race car driver kept track of how many laps he drove in the past 5 days. What is the mode of the numbers?

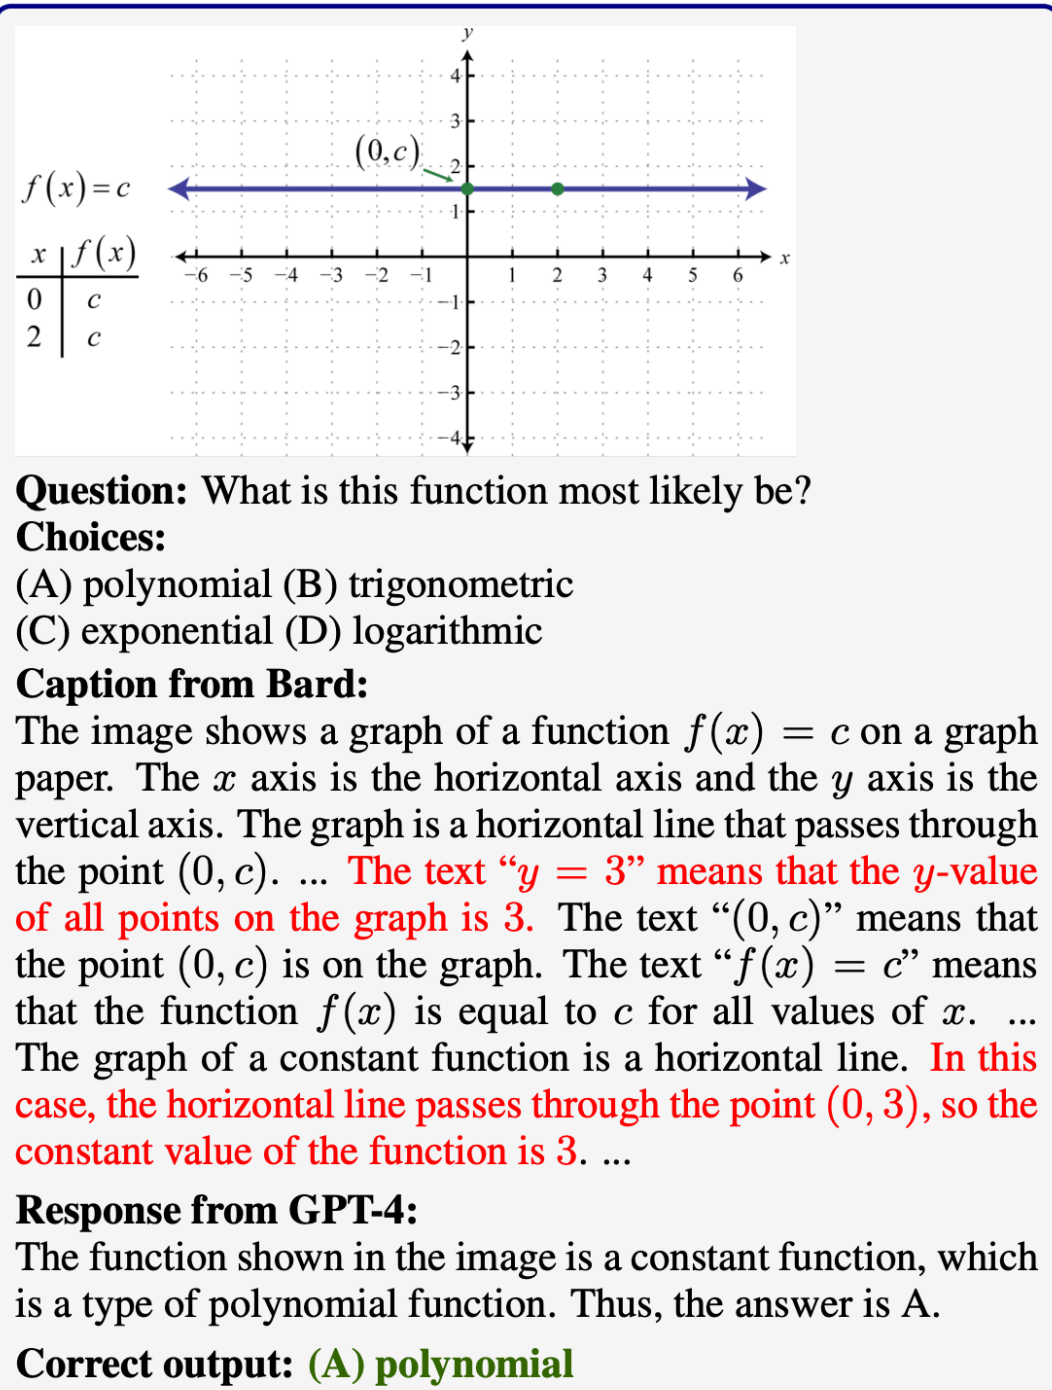
Detected OCR text:
(55, 9], Laps driven),
(21, 34], Day),
(81, 37], Number of laps),
(1, 64], Thursday), ([129, 67], 53),
(3, 99], Friday), ([129, 99], 53),
(1, 126], Saturday), ([129, 129], 53),
(3, 161], Sunday), ([129, 161], 55),
(1, 188], Monday), ([129, 191], 55)

Response from GPT-4:
def mode(laps):
 return max(set(laps), key=
 =laps.count)

laps = [53, 53, 53, 55, 55]
print(mode(laps))

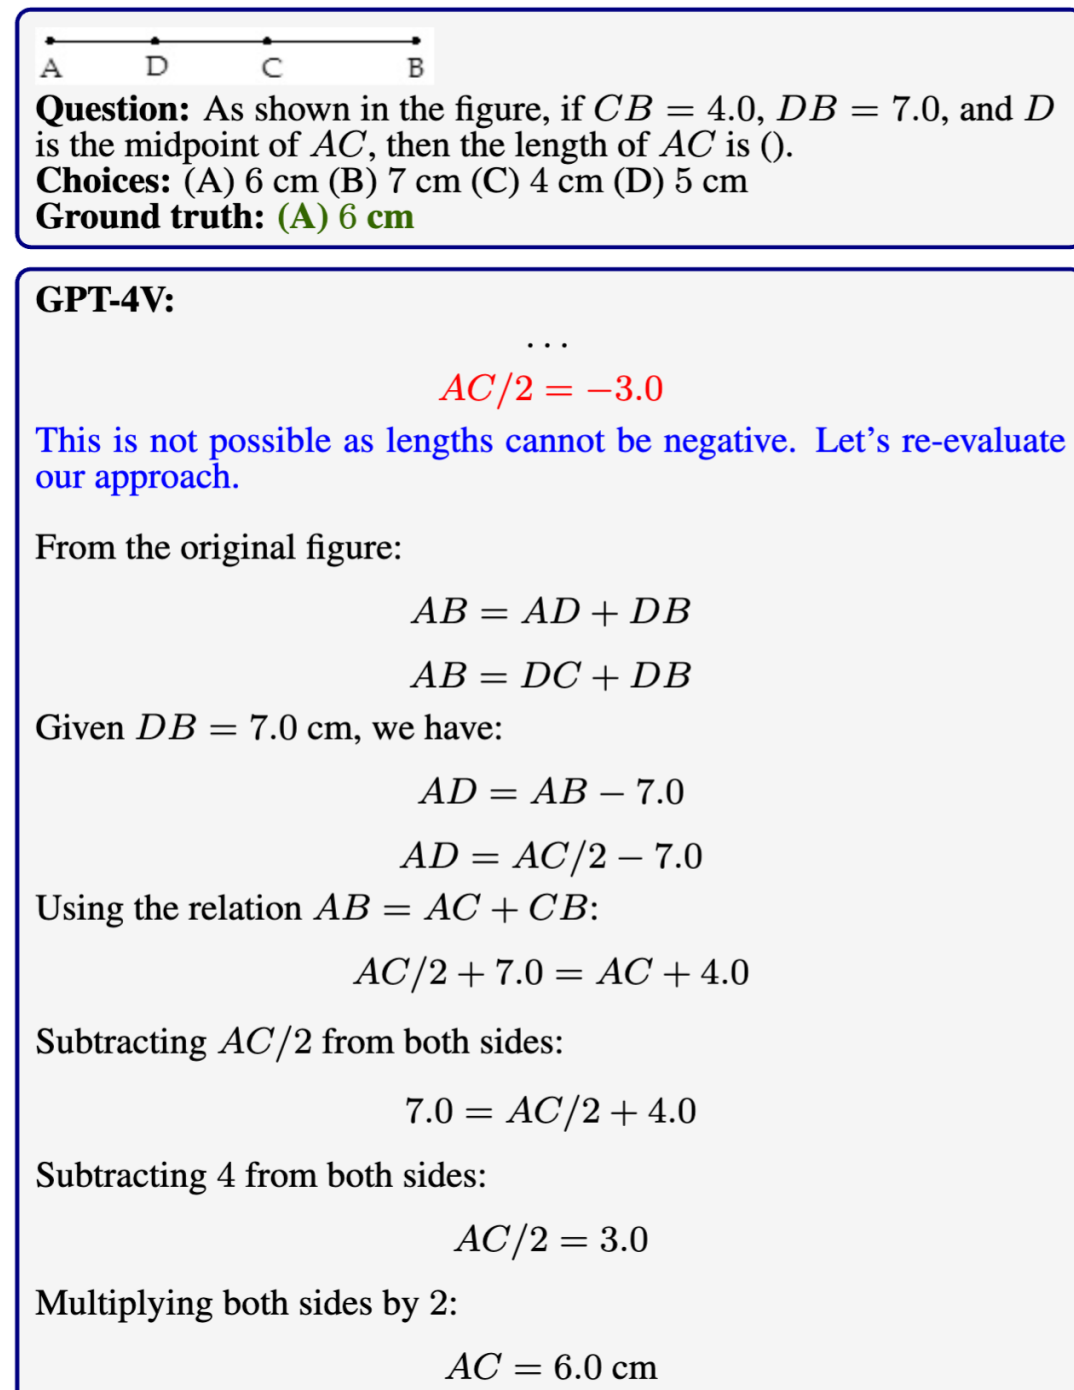
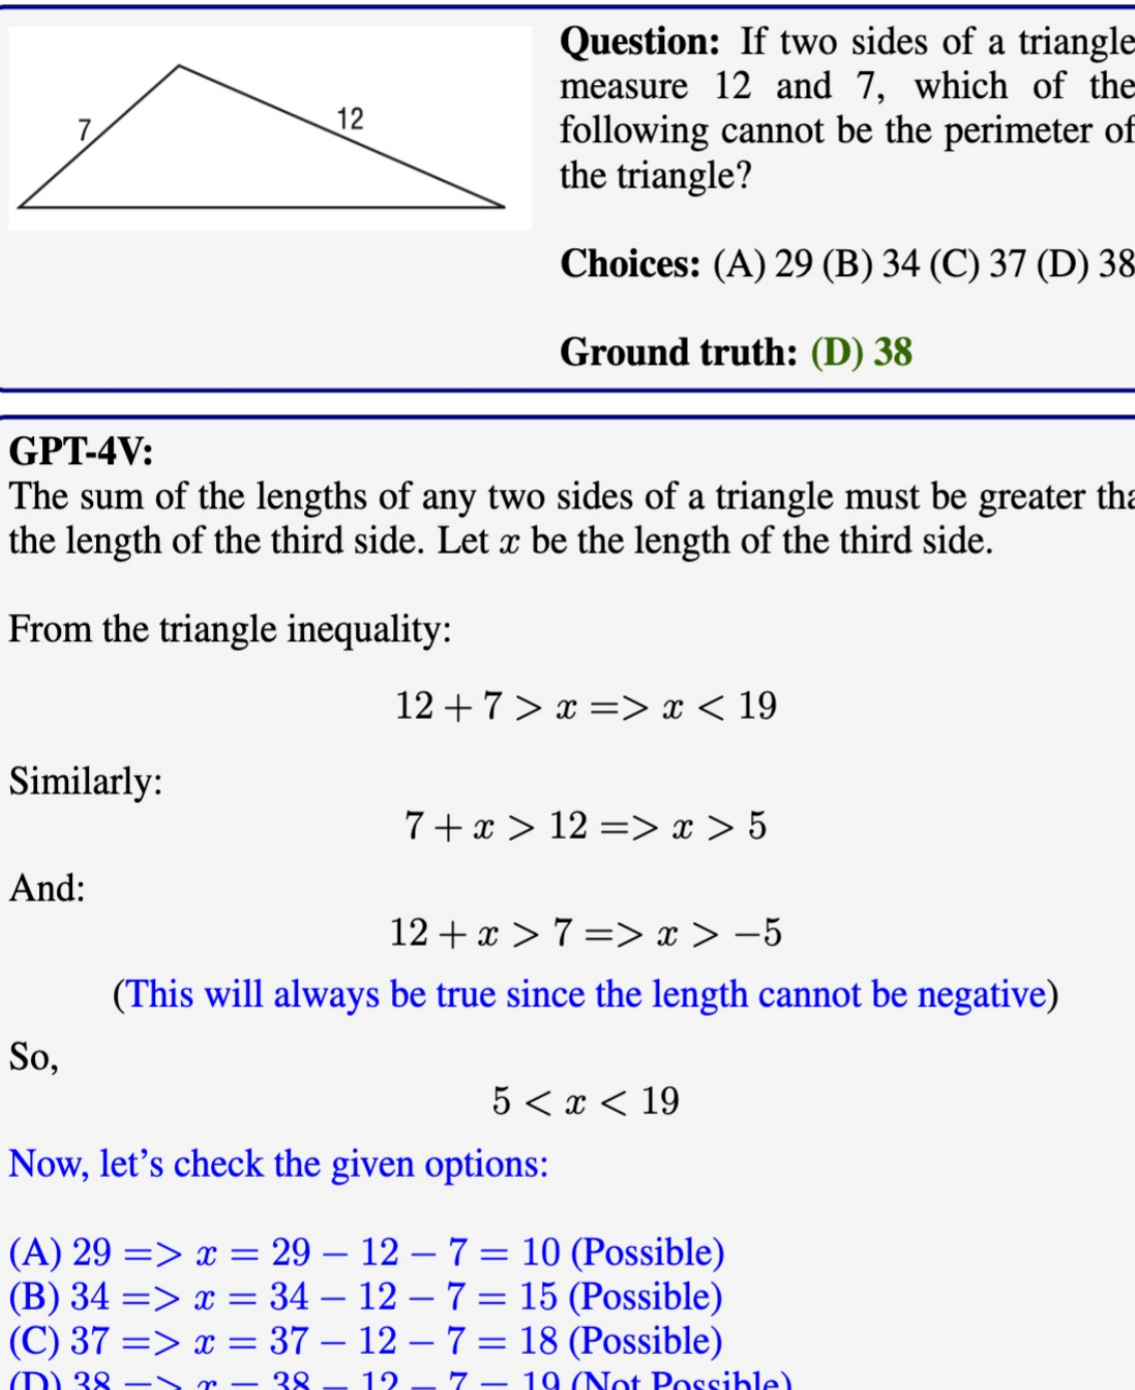
Correct output: 53

(a) Correct answer and code



(b) Correct answer with partially correct outputs

Self-Verification in GPT-4V



Self-Consistency for GPT-4V

Web pages visited	
Person	Number of Web pages
Bonnie	84
Winston	56
Percy	70
Derek	68
Jake	60
Finn	47
Carter	75
Kurt	52

ID: 331

Question: Several people compared how many Web pages they had visited. What is the mean of the numbers?

Ground truth: 64

GPT-4V (Reasoning Path 1):
To find the mean of the numbers, you add up all the numbers and then divide by the quantity of numbers.

Mean = $(84 + 56 + 70 + 68 + 60 + 47 + 75 + 52) \div 8$
 $= 462 \div 8$
 $= 57.75$

However, since the question requires an integer answer, the mean number of web pages visited is: 58.
Wrong output: 58

GPT-4V (Reasoning Path 2):
To find the mean number of web pages visited:

First, sum the number of web pages for all the individuals:
Bonnie: 84, Winston: 56, Percy: 70, Derek: 68, Jake: 60, Finn: 47, Carter: 75, Kurt: 52

Total = $84 + 56 + 70 + 68 + 60 + 47 + 75 + 52 = 512$

Next, divide by the number of individuals (8): Mean = $512 \div 8 = 64$

The mean number of web pages visited is 64.
Correct output: 64

GPT-4V (Reasoning Path 3):
First, I'll list the number of web pages each person visited:
Bonnie: 84, Winston: 56, Percy: 70, Derek: 68, Jake: 60, Finn: 47, Carter: 75, Kurt: 52

Next, I'll group some numbers together for easier addition:
 $(84 + 56) = 140$
 $(70 + 68) = 138$
 $(60 + 47) = 107$
 $(75 + 52) = 127$

Now, I'll sum these results: $140 + 138 + 107 + 127 = 512$

Now, I'll divide the sum by the number of individuals to get the mean: $512 \div 8 = 64$

So, the mean number of web pages visited is 64.
Correct output: 64

Multi-turn Human-AI Interaction in GPT-4V

